

COMP 251

Algorithms & Data Structures (Winter 2022)

Graphs – Single Source Shortest Paths

School of Computer Science
McGill University

Slides of (Comp321 ,2021), Langer (2014), Kleinberg & Tardos, 2005 & Cormen et al., 2009, Jaehyun Park' slides CS 97SI, Top-coder tutorials, T-414-AFLV Course, Programming Challenges books, slides from D. Plaisted (UNC) and Comp251-Fall McGill.

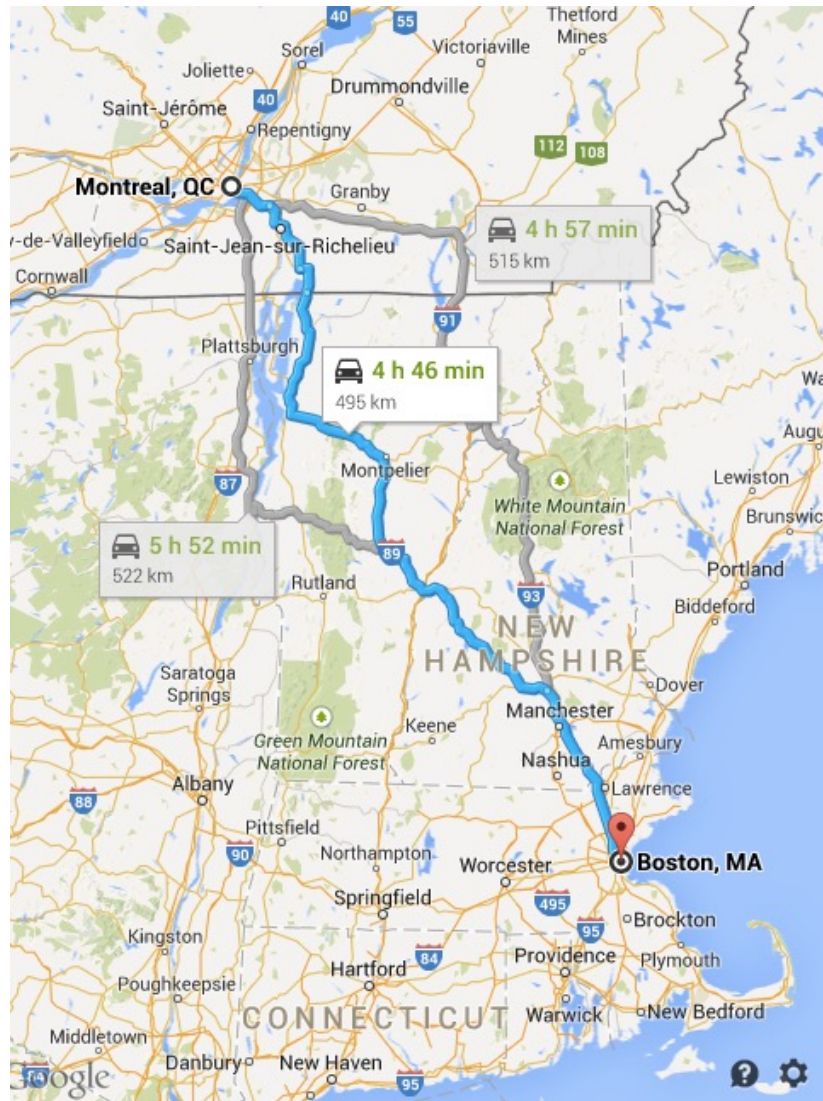
Announcements

- Deadline.
 - Peer review of 3 classmate submissions.

Outline

- Graphs.
 - Introduction.
 - Topological Sort. / Strong Connected Components
 - Network Flow 1.
 - Introduction
 - Ford-Fulkerson
 - Network Flow 2.
 - Min-cuts
 - Shortest Path.
 - Minimum Spanning Trees.
 - Bipartite Graphs.

Shortest Path - Problem



What is the shortest road to go from one city to another?

Example: Which road should I take to go from Montréal to Boston (MA)?

Variants:

- What is the fastest road?
- What is the cheapest road?

Shortest Path – As a graph problem

Input:

- Directed graph $G = (V, E)$
- Weight function $w: E \rightarrow \mathbb{R}$

Weight of path $p = \langle v_0, v_1, \dots, v_k \rangle$

$$= \sum_{k=1}^n w(v_{k-1}, v_k)$$

= sum of edges weights on path p


Shortest-path weight u to v :

$$\delta(u, v) = \begin{cases} \min \left\{ w(p) : u \xrightarrow{p} v \right\} & \text{If there exists a path } u \rightsquigarrow v. \\ \infty & \text{Otherwise.} \end{cases}$$

Shortest path u to v is any path p such that $w(p) = \delta(u, v)$


Generalization of breadth-first search to weighted graphs

Shortest Path – As a graph problem - example




Shortest path from s ?

Shortest Path – As a graph problem - example



Shortest paths are organized as a tree.
Vertices store the length of the shortest path from s .

Shortest Path – As a graph problem - example



*value needs
to be
minimal*

Shortest paths are not necessarily unique!

Shortest Path – As a graph problem - variants

- **Single-source (SSSP):** Find shortest paths from a given source vertex $s \in V$ to every vertex $v \in V$.
- **Single-destination:** Find shortest paths to a given destination vertex.
 - By reversing the direction of each edge in the graph, you reduce it to SSSP.
- **Single-pair:** Find shortest path from u to v .
 - SSSP solves this variant. All known algorithms for this problem have the same worst-case asymptotic running time as the best single-source algorithm.
- **All-pairs:** Find shortest path from u to v for all $u, v \in V$.
 - By running a SSSP algorithm once from each vertex, but we can solve it faster.



Shortest Path – Negative weight edges

Negative weight edges can create issues.

Why? If we have a negative-weight cycle, we can just keep going around it, and get $w(s, v) = -\infty$ for all v on the cycle.

When? If they are reachable from the source (Corollary: It is OK to have a negative-weight cycles if it is not reachable from the source).

What? Some algorithms work only if there are no negative-weight edges in the graph. We must specify when they are allowed and not.



Shortest Path – Cycles

Shortest paths cannot contain cycles:

- Negative-weight: Already ruled out.
- Positive-weight: we can get a shorter path by omitting the cycle.
- Zero-weight: no reason to use them \Rightarrow assume that our solutions will not use them.

Consequence:


- Since any acyclic path in a graph $G = (V, E)$ contains at most $|V|$ distinct vertices, it also contains at most $|V| - 1$ edges. Thus, we can restrict our attention to shortest paths of at most $|V| - 1$ edges.

Shortest Path – Optimal substructure

Lemma

Any subpath of a shortest path is a shortest path.

Proof: (cut and paste)



Suppose this path p is a shortest path from u to v .

Then $\delta(u, v) = w(p) = \underline{w(p_{ux})} + \underline{w(p_{xy})} + \underline{w(p_{yv})}$.


splitting
shortest path
in 3

Shortest Path – Optimal substructure

Lemma

Any subpath of a shortest path is a shortest path.

Proof: (cont'd)




Now suppose there exists a shorter path $x \overset{p'_{xy}}{\rightsquigarrow} y$.

Then $w(p'_{xy}) < w(p_{xy})$.

$$w(p') = w(p_{ux}) + w(p'_{xy}) + w(p_{yv}) < w(p_{ux}) + w(p_{xy}) + w(p_{yv}) = w(p).$$


Contradiction of the hypothesis that p is the shortest path!

Shortest Path – Customized BFS




Vertices count the number of edges used to reach them.


Shortest Path – Customized BFS




Shortest Path – Customized BFS



Shortest Path – Customized BFS



Shortest Path – Customized BFS



Can we generalize BFS to use edge weights?

SSSP – Output

For each vertex $v \in V$:

- $d[v] = \delta(s, v)$.
 - Initially, $d[v] = \infty$.
 - Reduces as algorithms progress, but always maintain $d[v] \geq \delta(s, v)$.
 - Call $d[v]$ a **shortest-path estimate**.
- $\pi[v] =$ predecessor of v on a shortest path from s .
 - If no predecessor, $\pi[v] = \text{NIL}$.
 - π induces a tree - **shortest-path tree** (see proof in textbook).

SSSP – Structure



1. Initialization

```
INIT-SINGLE-SOURCE( $V, s$ )  
  for each  $v \in V$  do  
     $d[v] \leftarrow \infty$  distances  
     $\pi[v] \leftarrow \text{NIL}$  predecessors  
   $d[s] \leftarrow 0$ 
```

SSSP – Structure

2. Scan vertices and relax edges.

- An edge $u \rightarrow v$ is tense if $d(u) + w(u, v) < d(v)$.
- If $u \rightarrow v$ is tense, the tentative shortest path $s \rightsquigarrow v$ is clearly incorrect, because the path $s \rightsquigarrow u \rightarrow v$ is shorter.



SSSP – Structure


2. Scan vertices and relax edges.

RELAX(u, v, w)



if $d[v] > d[u] + w(u, v)$ **then**

$d[v] \leftarrow d[u] + w(u, v)$

$\pi[v] \leftarrow u$



Relax



Shortest Path and Relaxation – Properties

- Triangle inequality
 - For any edge $(u, v) \in E$, we have $\delta(s, v) \leq \delta(s, u) + w(u, v)$.
- Upper-bound property
 - We always have $d[v] \geq \delta(s, v)$ for all vertices $v \in V$, and once $d[v]$ achieves the value $\delta(s, v)$, it never changes.
- No-path property
 - If there is no path from s to v , then we always have $d[v] = \delta(s, v) = \infty$.
- Convergence property.
 - If $s \rightsquigarrow u \rightarrow v$ is a shortest path in G for some $u, v \in V$, and if $d[u] = \delta(s, u)$ at any time prior to relaxing edge (u, v) , then $d[v] = \delta(s, v)$ at all time afterwards.
- Path-relaxation property
 - If $p = \langle v_0, v_i, \dots, v_k \rangle$ is a shortest path from $s = v_0$ to v_k , and we relax the edges of p in the order $(v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k)$ then $v_k[d] = \delta(s, v_k)$
- Predecessor-subgraph property
 - Once $v[d] = \delta(s, v)$ for all $v \in V$, the predecessor subgraph is a shortest-paths tree rooted at s . (seen before)

SSSP – DAG

Directed acyclic graph
(no neg cycles)

DAG \Rightarrow no negative-weight cycles.

DAG-SHORTEST-PATHS (V, E, w, s)


topologically sort the vertices

INIT-SINGLE-SOURCE (V, s)

for each vertex u in topological order **do**


for each vertex $v \in \text{Adj}[u]$ **do**

RELAX(u, v, w)




SSSP – DAG - Example

topologically sorted graph




SSSP – DAG - Example

relax edges of S




SSSP – DAG - Example


(xv4.) relax edges from t




SSSP – DAG - Example



SSSP – DAG - Example



SSSP – DAG - Example



SSSP – DAG

DAG-SHORTEST-PATHS (V, E, w, s)

topologically sort the vertices $\longrightarrow O(V+E)$

INIT-SINGLE-SOURCE (V, s) $\longrightarrow O(V)$

for each vertex u in topological order **do**

for each vertex $v \in Adj[u]$ **do**

 RELAX (u, v, w) \leftarrow constant time $O(1)$

$O(V+E)$

\uparrow
graph
traversal

Time: $(V + E)$.

Correctness:

Because we process vertices in topologically sorted order,
edges of any path must be relaxed in order of appearance in the path. \Rightarrow Edges on any shortest path are relaxed in order.
 \Rightarrow By path-relaxation property, correct.

SSSP – Dijkstra's algorithm

- No negative-weight edges.
- Weighted version of BFS:
 - **Instead of a FIFO queue, uses a priority queue.**
 - Keys are shortest-path weights ($d[v]$).
- Have two sets of vertices:
 - S = vertices whose final shortest-path weights are determined,
 - Q = priority queue = $V - S$.
- Greedy choice: At each step we choose the light edge.

BFS w/
priority Q

SSSP – Dijkstra's algorithm

DIJKSTRA(V, E, w, s)

INIT-SINGLE-SOURCE(V, s) *initialization*

$S \leftarrow \emptyset$ *← set of determined shortest-path weight*

$Q \leftarrow V$ *← vertices that still need to be processed*

Q not empty [**while** $Q \neq \emptyset$ **do** *process the Q*

$u \leftarrow \text{EXTRACT-MIN}(Q)$ *Greedy choice*


$S \leftarrow S \cup \{u\}$ *union/append u to S*

for each vertex $v \in \text{Adj}[u]$ **do**

RELAX(u, v, w)

has an edge with u

SSSP – Dijkstra's algorithm - Example




initialization
sets
dist $s \rightarrow$ node
to infinity

min dist is 0

Q


s	t	y	x	z
0	∞	∞	∞	∞

SSSP – Dijkstra's algorithm - Example




SSSP – Dijkstra's algorithm - Example

Process 4




relax
orange
edges



SSSP – Dijkstra's algorithm - Example

process t




relax yellow edges




SSSP – Dijkstra's algorithm - Example

process x




SSSP – Dijkstra's algorithm - Example


process 2



Q



SSSP – Dijkstra's algorithm - Example



Dijkstra's algorithm - Correctness

Loop invariant:

At the start of each iteration of the while loop, $d[v] = \delta(s, v)$ for all $v \in S$.

Initialization:

Initially, $S = \emptyset$, so trivially true.

Termination:

At end, $Q = \emptyset \Rightarrow S = V \Rightarrow d[v] = \delta(s, v)$ for all $v \in V$.

Maintenance:

Show that $d[u] = \delta(s, u)$ when u is added to S in each iteration.

while $Q \neq \emptyset$

distance estimate of v

shortest path

once v is added to S

Dijkstra's algorithm - Correctness


Maintenance

Show that $d[u] = \delta(s, u)$ when u is added to S in each iteration.

Suppose there exists u such that $d[u] \neq \delta(s, u)$.

Let u be the first vertex for which $d[u] \neq \delta(s, u)$ when u is added to S .

- $u \neq s$, since $d[s] = \delta(s, s) = 0$.
- Therefore, $s \in S$, so $S \neq \emptyset$.
- There must be some path $s \rightsquigarrow u$. Otherwise $d[u] = \delta(s, u) = \infty$ by no-path property.
- So, there is a path $s \rightsquigarrow u$. Thus, there is a shortest p path $s \rightsquigarrow u$.




Dijkstra's algorithm - Correctness

Show that $d[u] = \delta(s, u)$ when u is added to S in each iteration.

Just before u is added to S , the path p connects a vertex in S (i.e., s) to a vertex in $V - S$ (i.e., u).

Let y be the first vertex along p that is in $V - S$ and let $x \in S$ be the predecessor of y .



Decompose p into $s \xrightarrow{p_1} x \rightarrow y \xrightarrow{p_2} u$.


Dijkstra's algorithm - Correctness

Claim 1: $d[y] = \delta(s, y)$ when u is added to S .

Proof:

$x \in S$ and u is the first vertex such that $d[u] \neq \delta(s, u)$ when u is added to S
 $\Rightarrow d[x] = \delta(s, x)$ when x is added to S .

But when x is added we relax the edge (x, y) , so by the convergence property, $d[y] = \delta(s, y)$.



Dijkstra's algorithm - Correctness

Show that $d[u] = \delta(s, u)$ when u is added to S in each iteration.

Now, we can get a contradiction to $d[u] \neq \delta(s, u)$:

y is on shortest path $p(s \rightsquigarrow u)$, and all edge weights are nonnegative.

$\Rightarrow \delta(s, y) \leq \delta(s, u)$ (y occurs before u in shortest path)

Then by Claim 1, we have $d[y] = \delta(s, y)$

$$\leq \delta(s, u)$$

$$\Rightarrow \leq d[u] \quad (\text{upper-bound property})$$

In addition, **y and u were in Q when we chose u , thus $d[u] \leq d[y]$.**

We have $d[y] \leq d[u]$ & $d[u] \leq d[y] \Rightarrow d[u] = d[y]$.

Therefore, $d[y] = \delta(s, y) \leq \delta(s, u) \leq d[u] = d[y]$

Contradicts assumption that $d[u] \neq \delta(s, u)$. ■

Dijkstra's algorithm - Analysis

DIJKSTRA(V, E, w, s)

INIT-SINGLE-SOURCE(V, s)

$S \leftarrow \emptyset$

$Q \leftarrow V$

while $Q \neq \emptyset$ **do**

$u \leftarrow \text{EXTRACT-MIN}(Q)$

$S \leftarrow S \cup \{u\}$

for each vertex $v \in \text{Adj}[u]$ **do**

 RELAX(u, v, w)

INSERT $O(?)$

$O(V)$

EXTRACT_MIN
 $O(?)$

$O(E)$

DECREASE-KEY
 $O(?)$

Dijkstra's algorithm - Analysis

It depends on implementation of priority queue.

If binary heap, each operation takes $O(\lg V)$ time
 $\Rightarrow O(E \lg V)$.

Note: We can achieve $O(V \lg V + E)$ with Fibonacci heaps.

Outline

- Graphs.
 - Introduction.
 - Topological Sort. / Strong Connected Components
 - Network Flow 1.
 - Introduction
 - Ford-Fulkerson
 - Network Flow 2.
 - Min-cuts
 - Shortest Path.
 - Minimum Spanning Trees.
 - Bipartite Graphs.

